

# Cloud-Hydra: A Cloud Native Multi-Cloud Defensive Load Balancing !Framework

Josh Stern, Rachid Tak Tak, Julian Trinh, Filip Vukelic

**SPRINT 2**

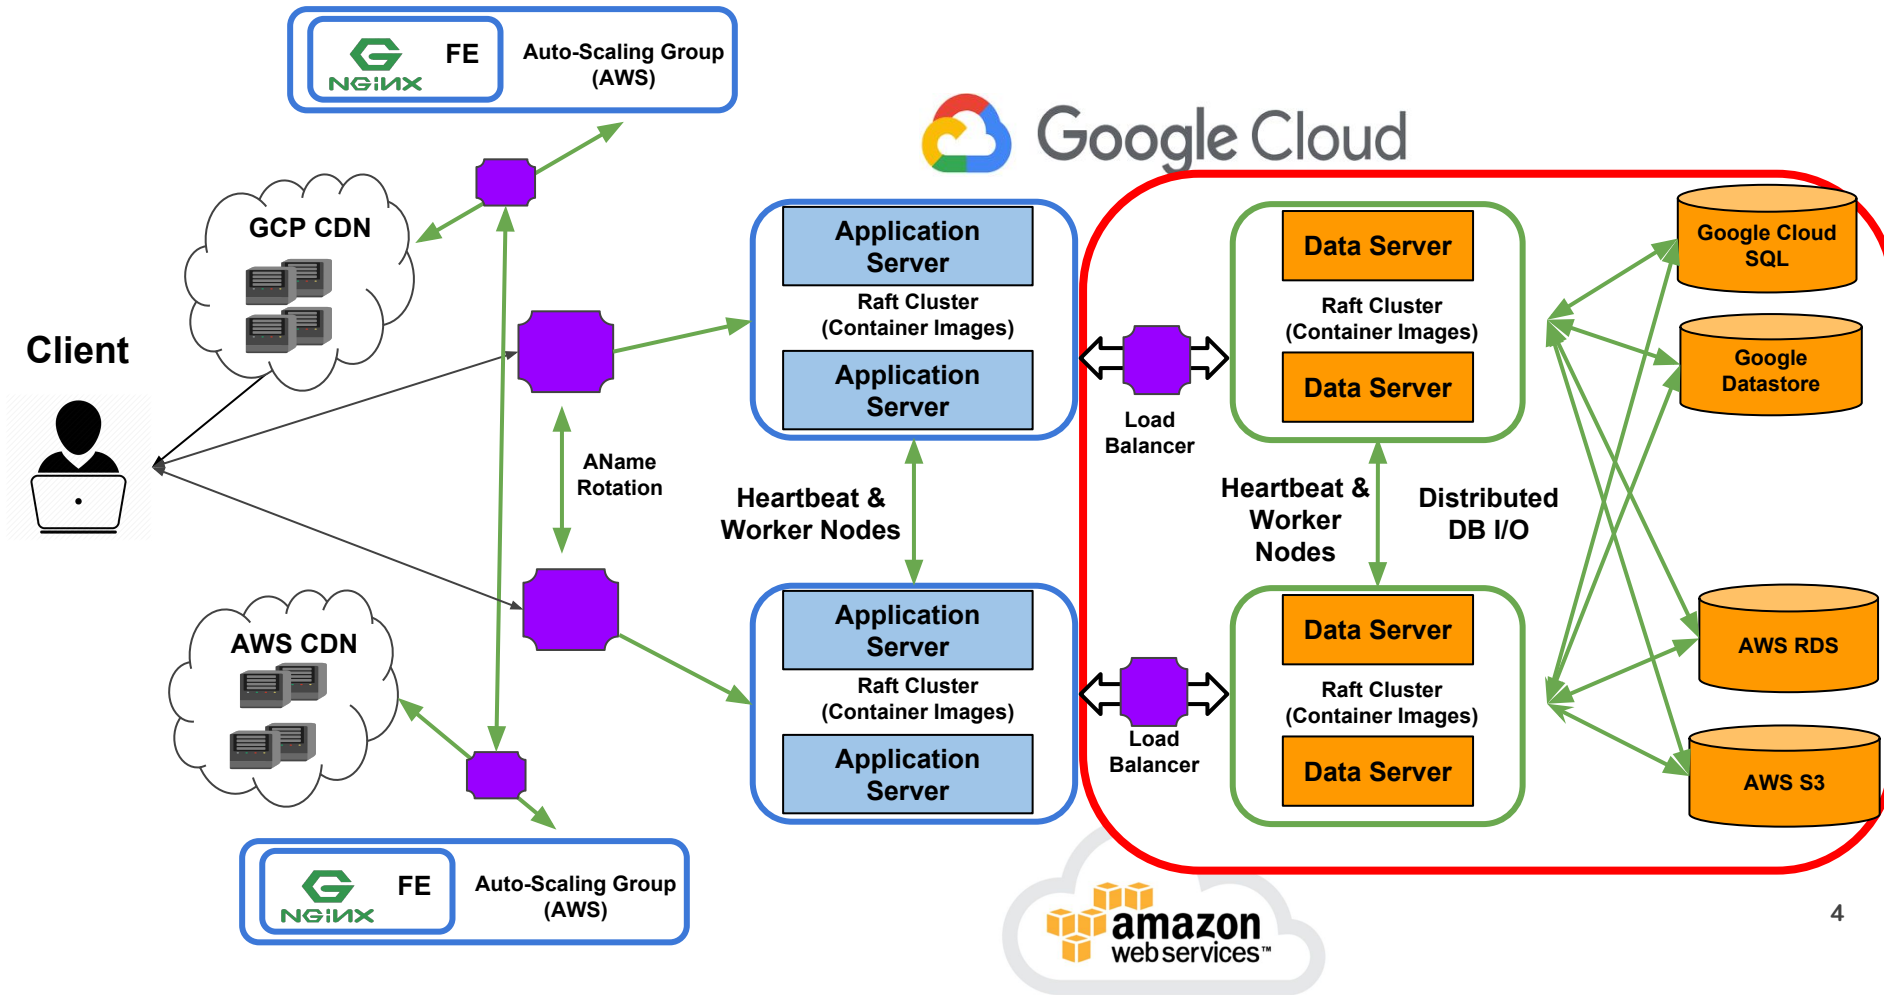


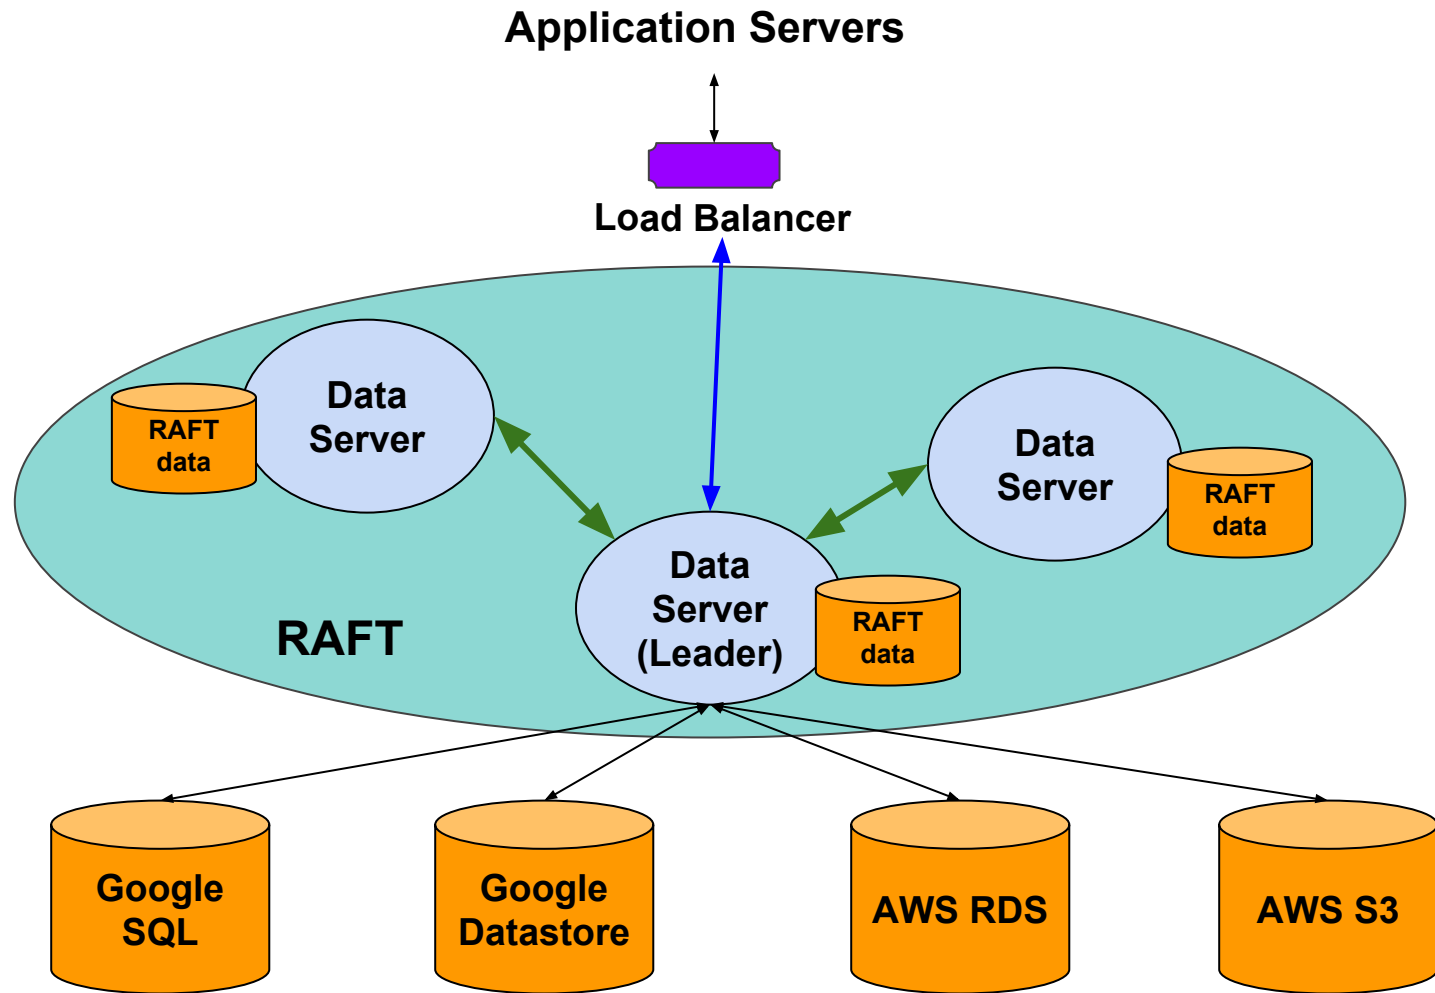
# Goals of Sprint 2

- Mainly focused on the the data layer
- Load balance requests from the application servers
- Unified data layer that abstracts individual servers
- Data servers configured in a Raft cluster

# Consensus (With Raft)

- Agreement on a unified cluster state by the cluster nodes
- Why? Prevent duplicate writes, defending against failover
- Using Raft -> simpler than Paxos, Good Hashicorp implementation in GO which was simple to modify
- Persistent store on each node with leader for state agreement (Term/Round, latest transaction counter and majority of nodes have it)
- How do we know which is accurate data? Leader or if no leader: term, latest transaction counter
- DBs can go down and need to be caught up



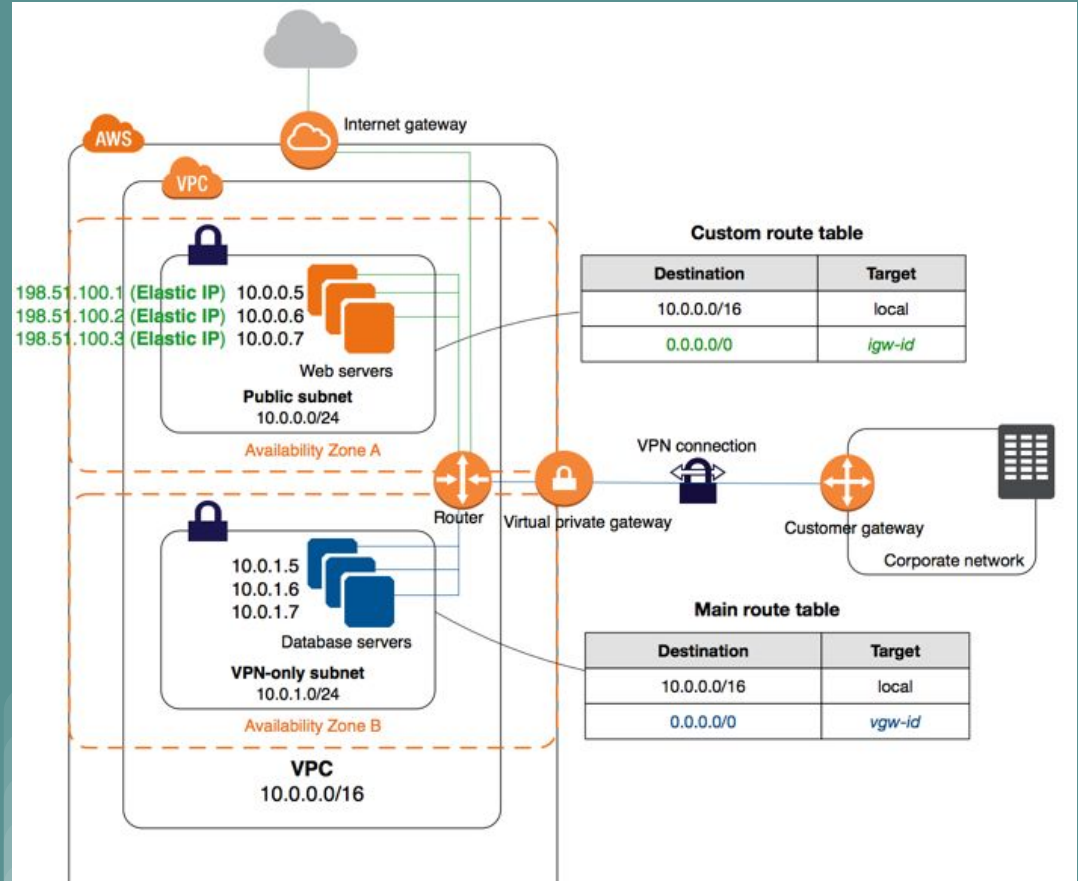


# Data Layer Raft Cluster

- Assume fail-stop model
- To handle  $n$  failures, need  $2n+1$  nodes
- Leader of the cluster handles reads and writes
- Writes are committed to all nodes via consensus before an actual DB write happens
- GETs (theoretically) may be serviced by any node
- When a node comes back up, writes that have occurred will be forwarded to it
- New leader will be elected if a leader dies

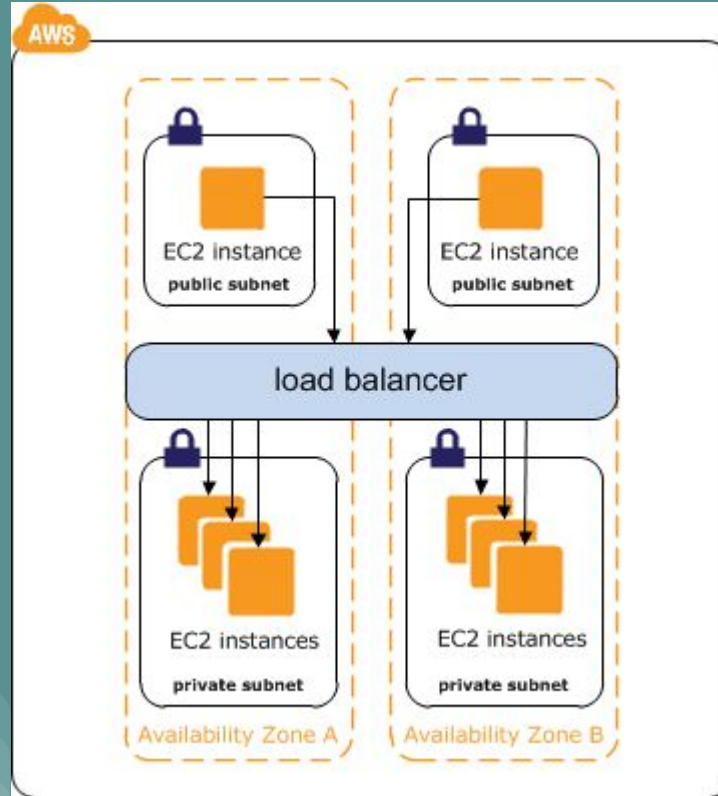
# Cross Cloud Consensus

- Want to keep data servers within private VPC (Virtual Private Cloud)
- Need to be able to communicate to cross cloud VPCs using VPN (Virtual Private Network)



# Cross Cloud Load Balancing

- Tried to do cross cloud load balancing
- AWS makes it very hard
- To load balance a private subnet, the load balancer must be private
- Heartbeat problems
- Easier solution: Nginx





# Single Cloud Load Balancing

- Load balance across each cloud separately using Nginx
- Use cross cloud VPN communication for (only) consensus on each layer
- Already implemented Nginx load balancing to dao servers which have consensus
  - Still needed: leader forwarding and replicating this on each layer.



# DEMO TIME!

# Sprint 2 Burndown



100% ∨ 79 total points


79 completed points

0 open tasks

22 closed tasks



0 cocaine doses



<https://tree.taiga.io/project/bowenislandsong-multi-cloud-defensive-load-balancing/taskboard/sprint-2-7327>

# Next Steps

- Raft leader forwarding
- Detecting when databases are down
- Forwarding recovered databases to the current state
- Load balancing traffic from the application server to cross-cloud data layer

**Thank you!**  
**Questions?**