

# Cloud-Hydra: A Cloud Native Multi-Cloud Defensive Load Balancing Framework

Josh Stern, Rachid Tak Tak, Julian Trinh, Filip Vukelic

**SPRINT 4**

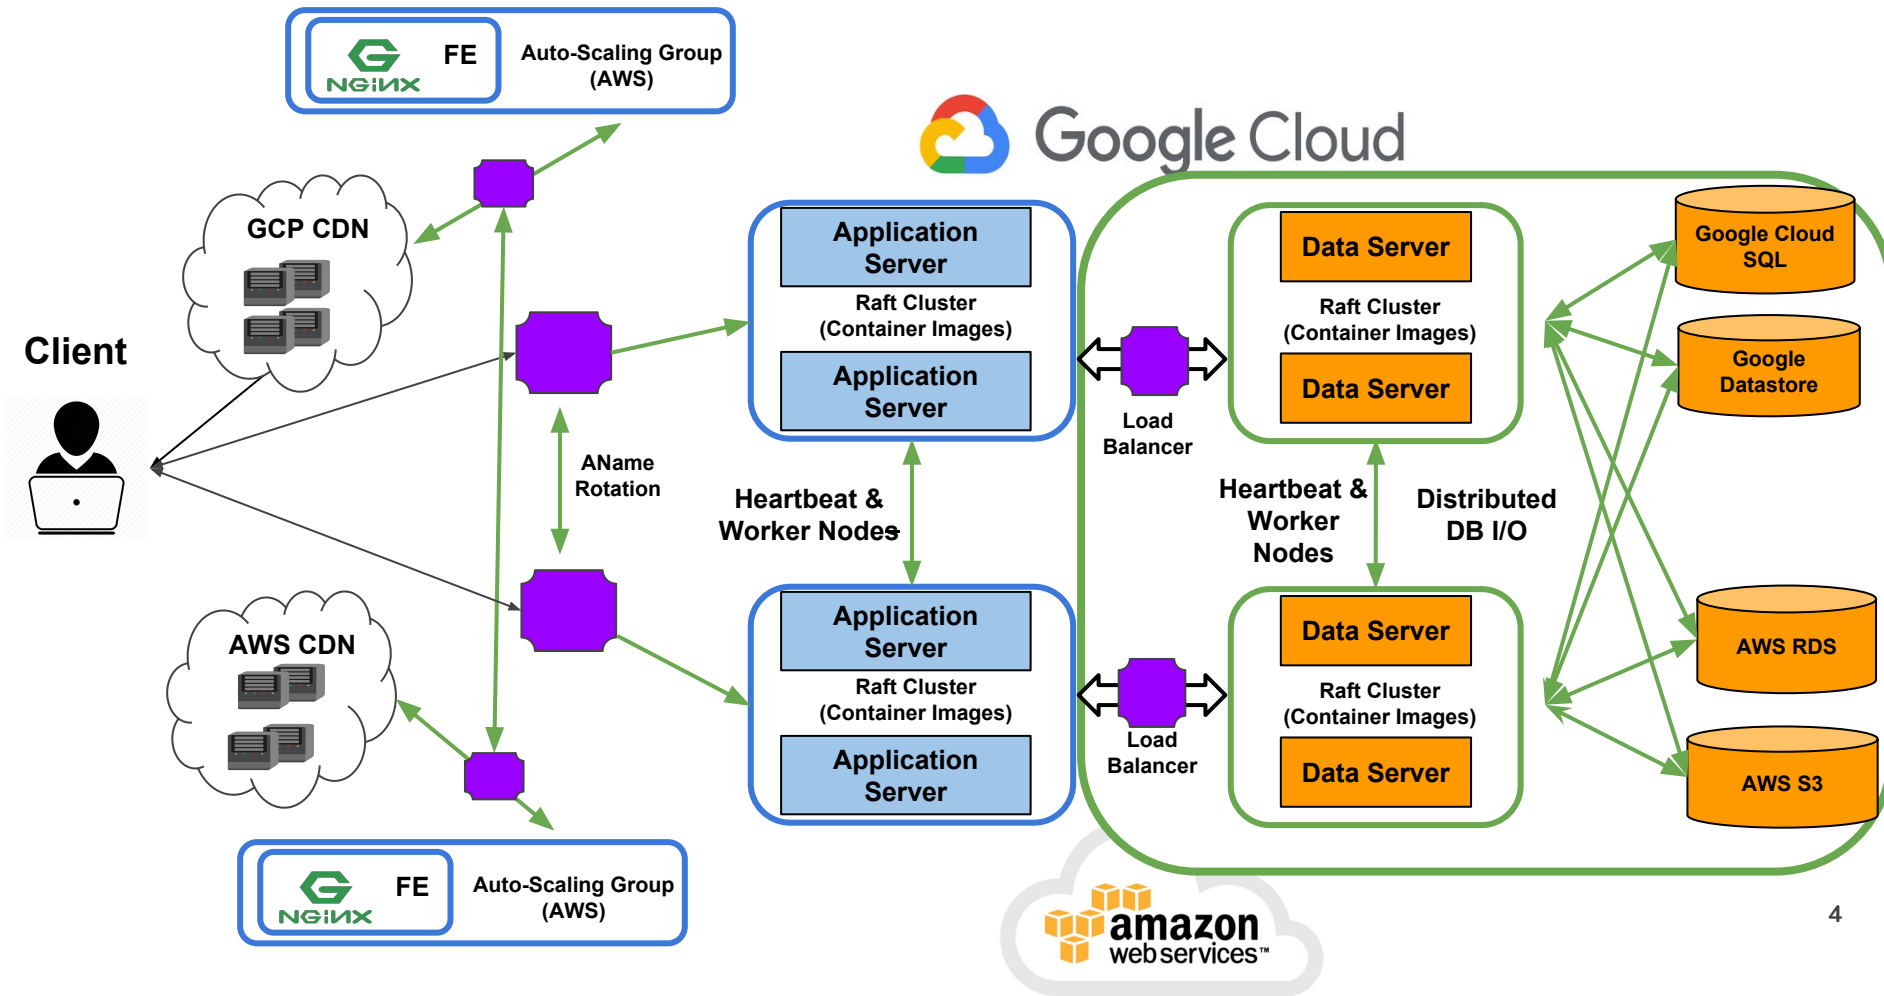


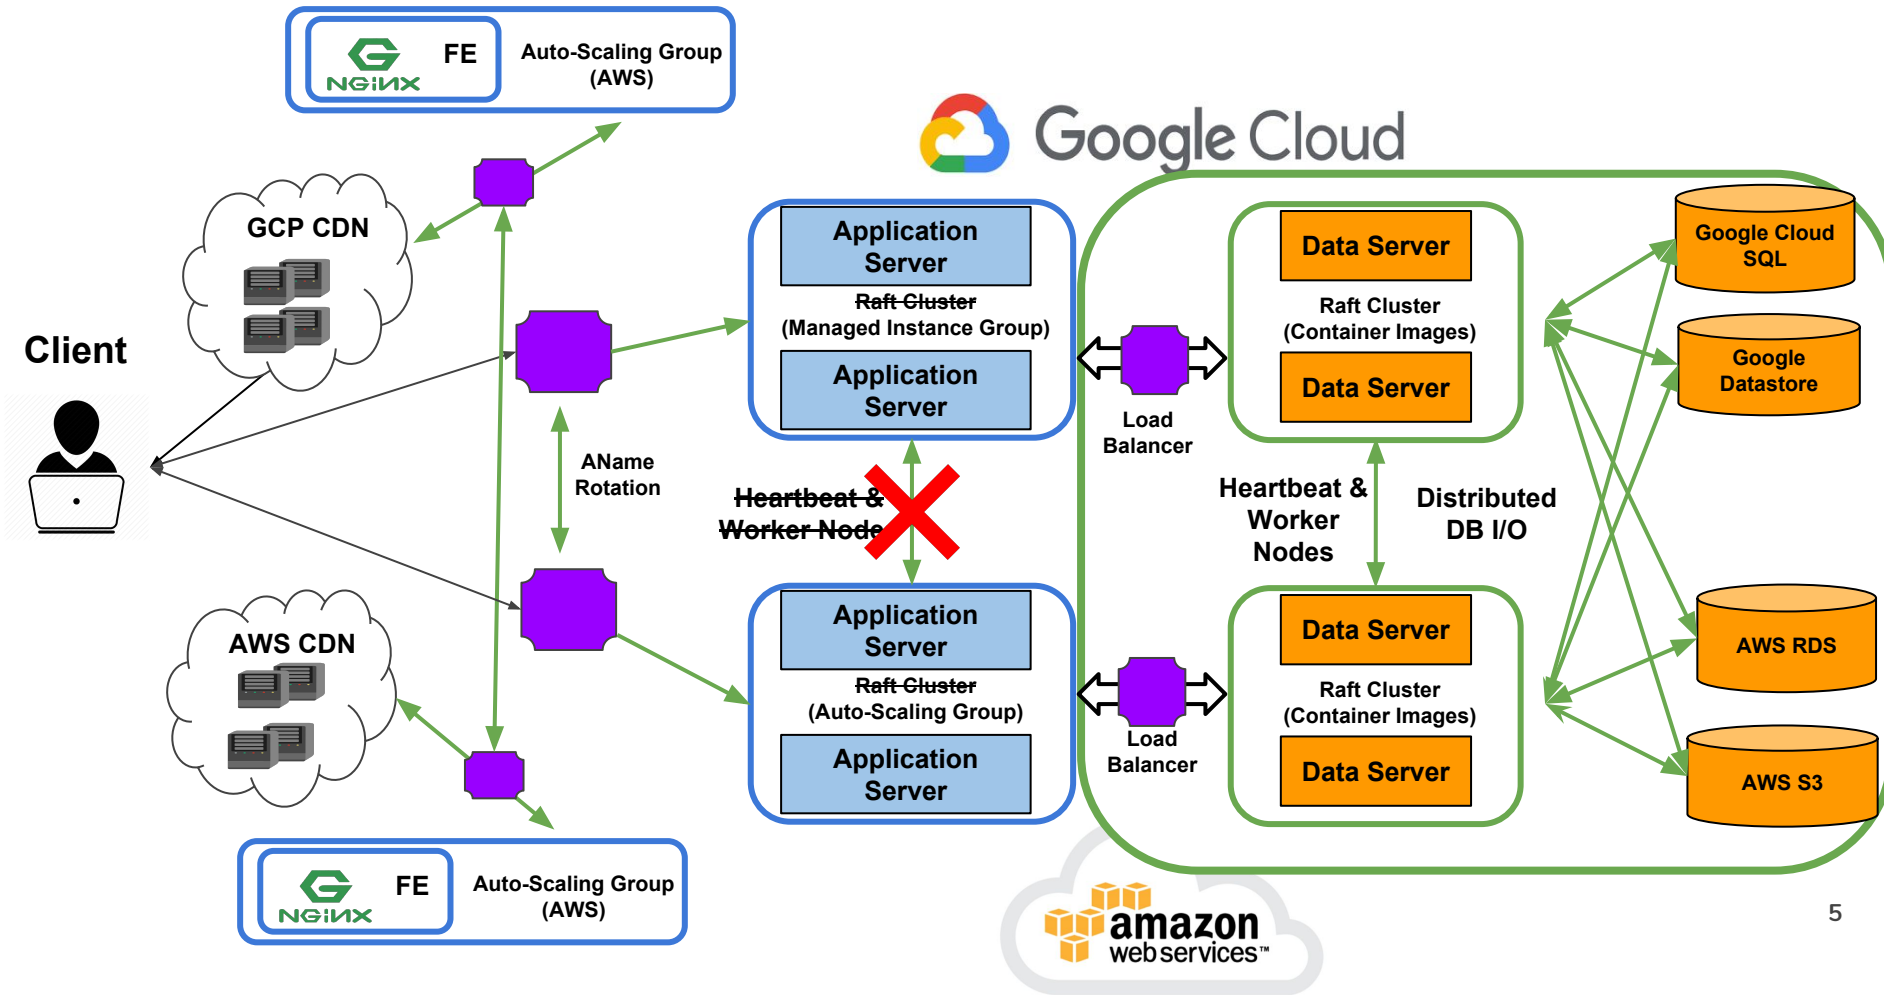
# Sprint History

- **Sprint 1- Built fullstack web app in GCP and AWS**
  - Compute Engine
  - Managed Instance Groups (with Docker Images)
  - Nginx front end server + CDNs
- **Sprint 2 - Unified Data Layer**
  - Raft consensus
  - New leader election
  - Distributed Raft log
  - Cross-cloud consensus, single cloud load balancing
- **Sprint 3 - Cross cloud forwarding and recovery**
  - Leader Forwarding
  - Data layer Load Balancing
  - Database Recovery

# Goals of Sprint 4

- Testing Raft Cluster I/O with node failures
- Testing Multi-Cloud DBRecovery
- Synchronizing AWS and GCP databases



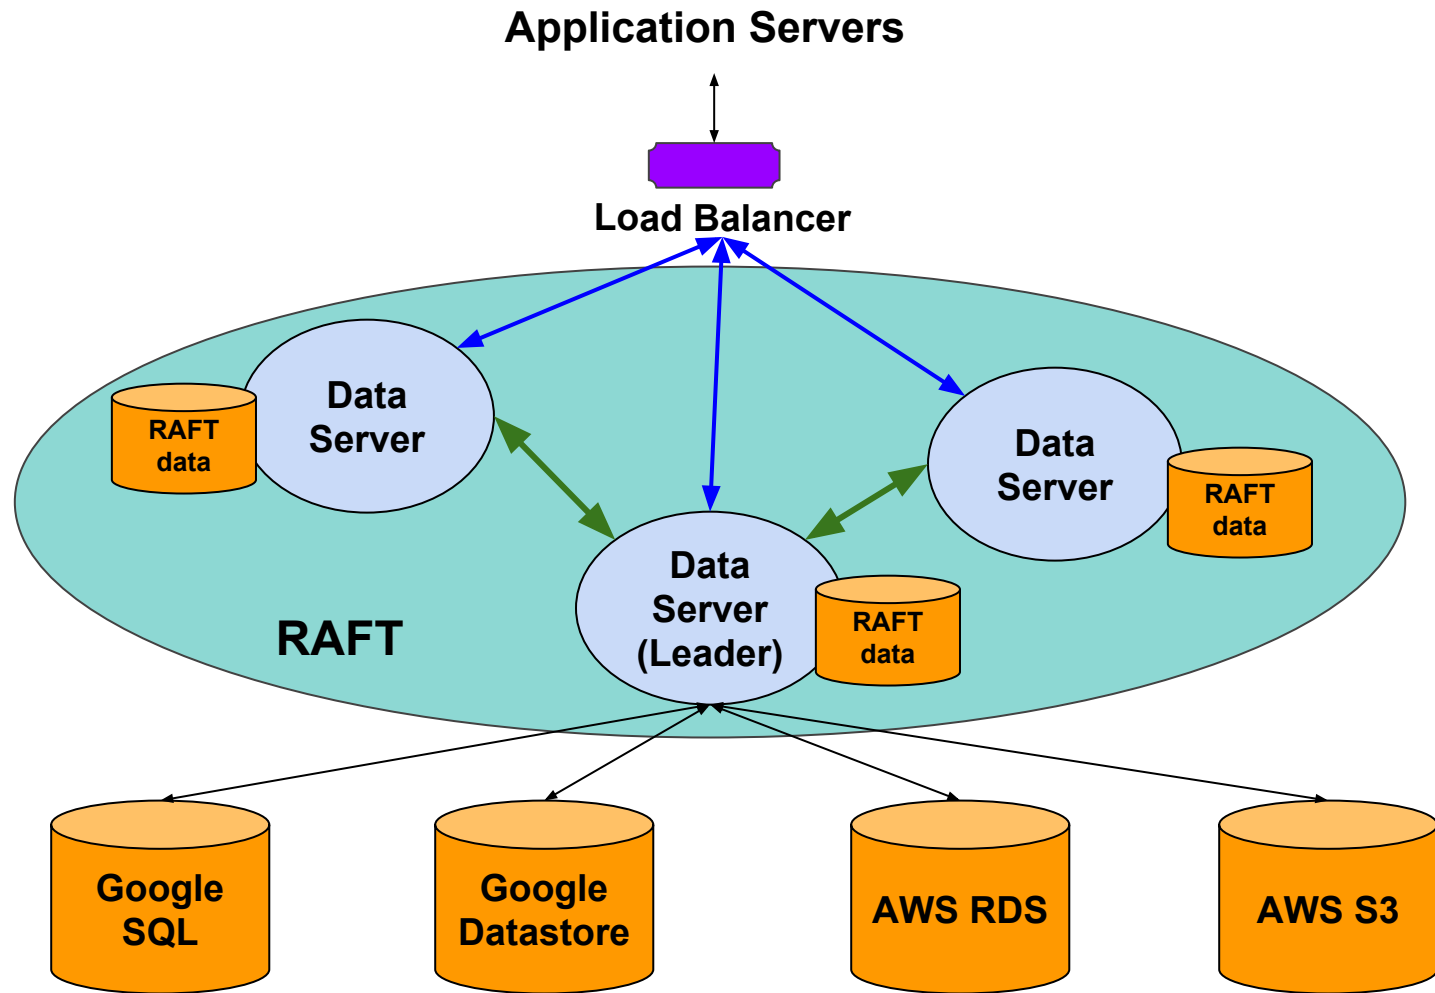


# Application Layer is Stateless

- Originally planned to configure the application layer into a Raft cluster
- The application (webserver) doesn't care about state
- Can be treated like a microservice to scale horizontally without much overhead

# Configuring the Data Layer Load Balancer

- Using Nginx as a proxy
- Requests are distributed to the cluster via Round Robin
- Nginx will periodically heartbeat with cluster nodes to ensure availability before scheduling a request





**DEMO TIME !**

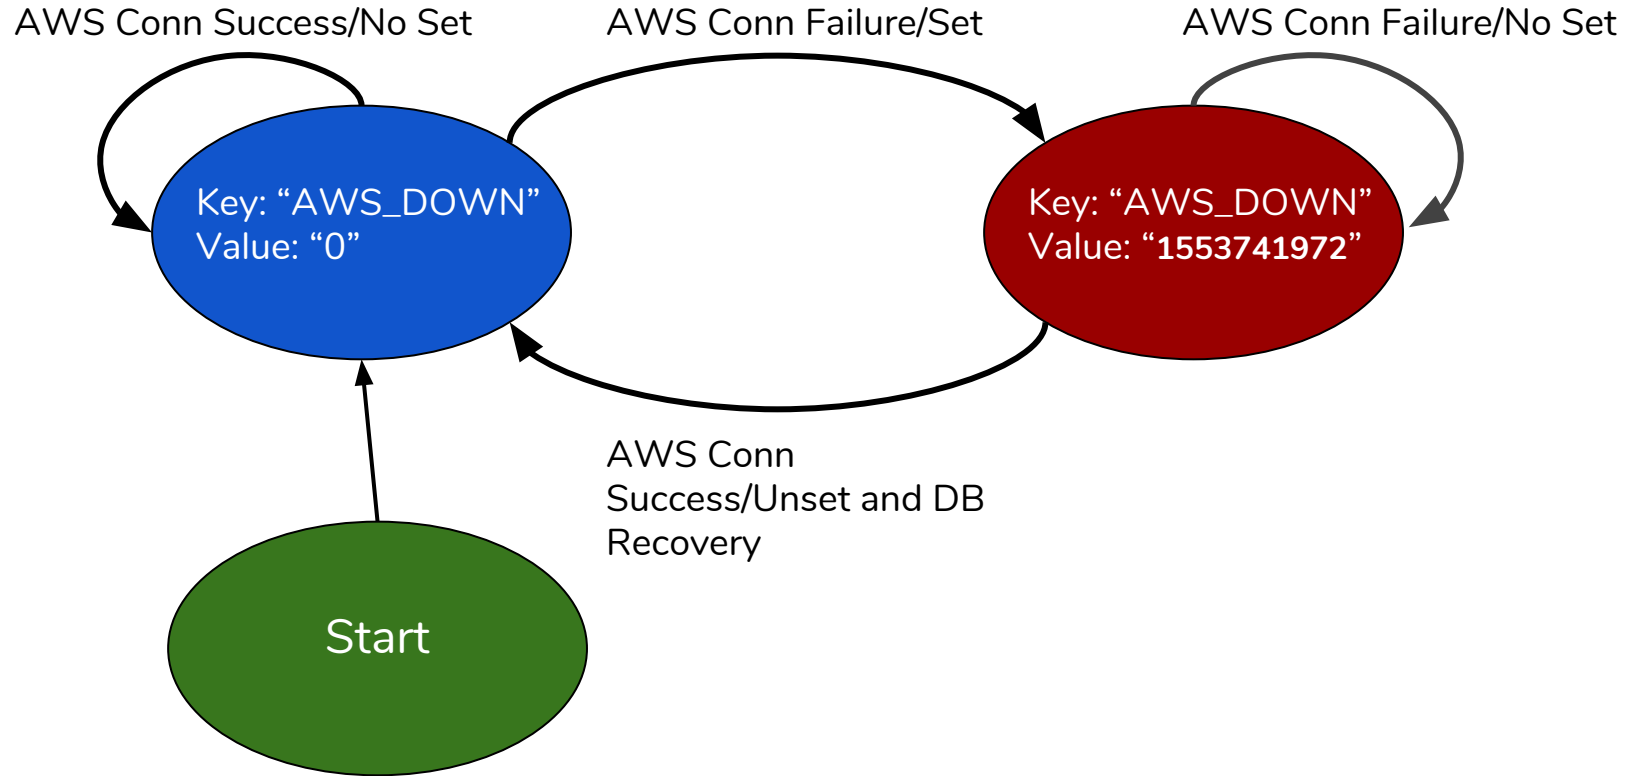
# Multi-Cloud Database I/O

- The database servers will read from a pseudo random database that is healthy
- If read encounters unhealthy DB, it will notify leader and read from a healthy DB
- On writes, the dao servers will write to all the DBs
- If one or all the dbs are down, beside the instructions, the db down state will be written to the raft log as well

# Database Forwarding


- Detecting down databases
- Detecting databases that have come back from the dead
  - Use raft log for recovery
- If a non leader node detects a database state change, notify the leader
  - Only the leader can propose changes to the raft log / update db state
- Internal Timestamps -- need to switch to external source

# Database Recovery



# Database Recovery Example

2:00 pm: Client issues write (Insert Warren Towers into Garages)



**Raft log:**

{GCP\_DOWN, 0}


{AWS\_DOWN, 0}

{Insert Warren Towers into Garages, timestamp: 2:00pm}

# Database Recovery Example

2:00 pm: Client issues write (Insert Warren Towers into Garages)

2:05 pm: AWS RDS goes down



**Raft log:**

{GCP\_DOWN, 0}

{AWS\_DOWN, 0}

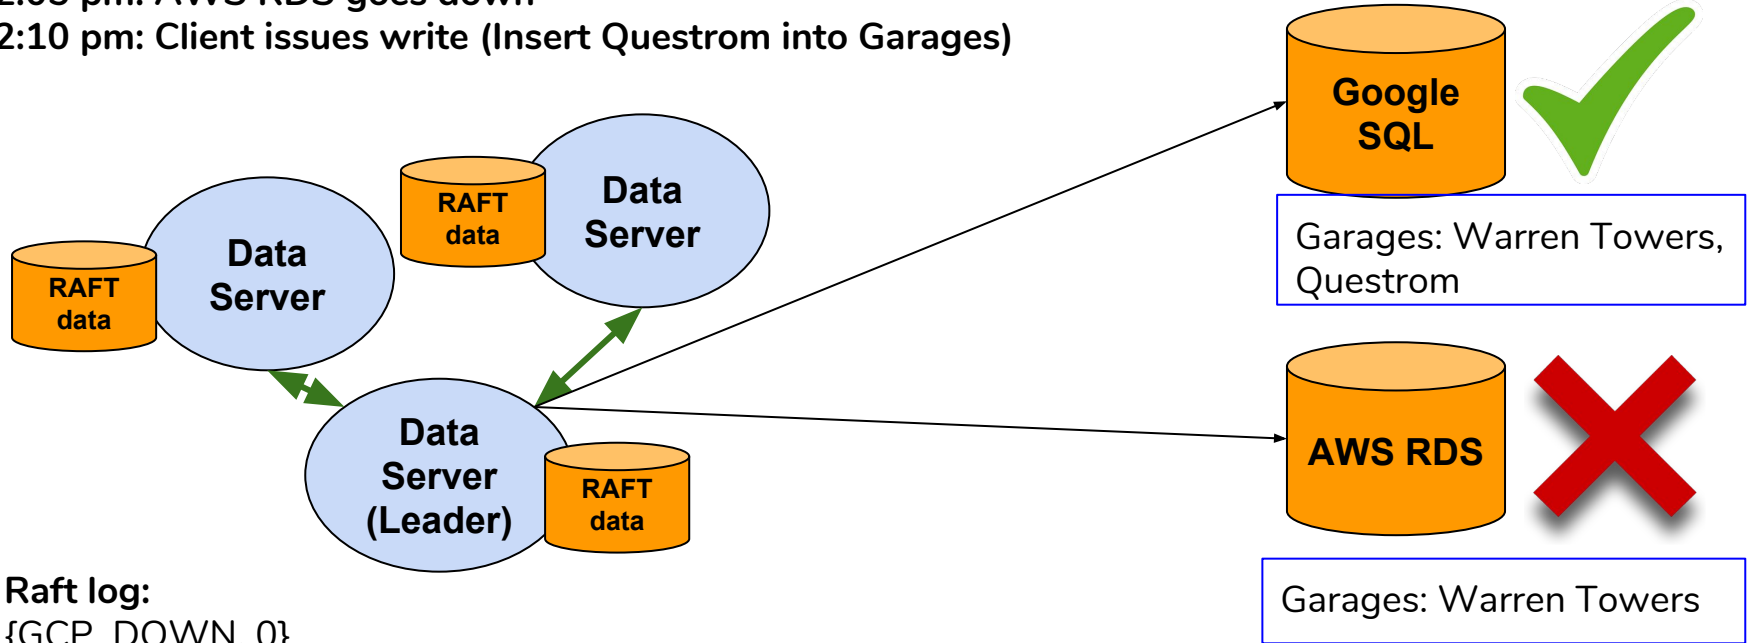
{Insert Warren Towers into Garages, timestamp: 2:00pm}

# Database Recovery Example

2:00 pm: Client issues write (Insert Warren Towers into Garages)

2:05 pm: AWS RDS goes down

2:10 pm: Client issues write (Insert Questrom into Garages)



## Raft log:

{GCP\_DOWN, 0}

{AWS\_DOWN, 2:10pm}

{Insert Warren Towers into Garages, timestamp: 2:00pm}

{Insert Questrom into Garages, timestamp: 2:10pm}

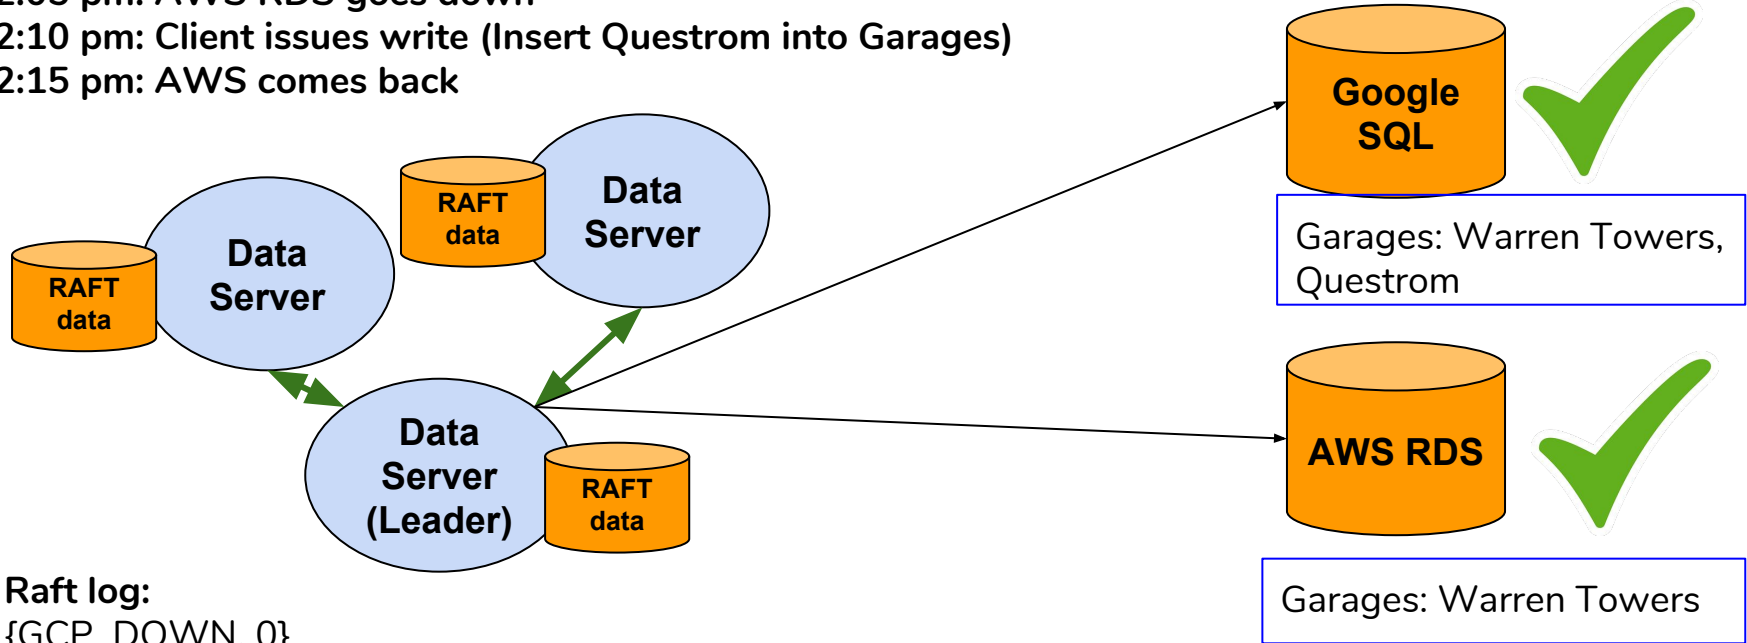
# Database Recovery Example

2:00 pm: Client issues write (Insert Warren Towers into Garages)

2:05 pm: AWS RDS goes down

2:10 pm: Client issues write (Insert Questrom into Garages)

2:15 pm: AWS comes back



## Raft log:

```
{GCP_DOWN, 0}  
{AWS_DOWN, 2:10pm}  
{Insert Warren Towers into Garages, timestamp: 2:00pm}  
{Insert Questrom into Garages, timestamp: 2:10pm}
```



# Database Recovery Example

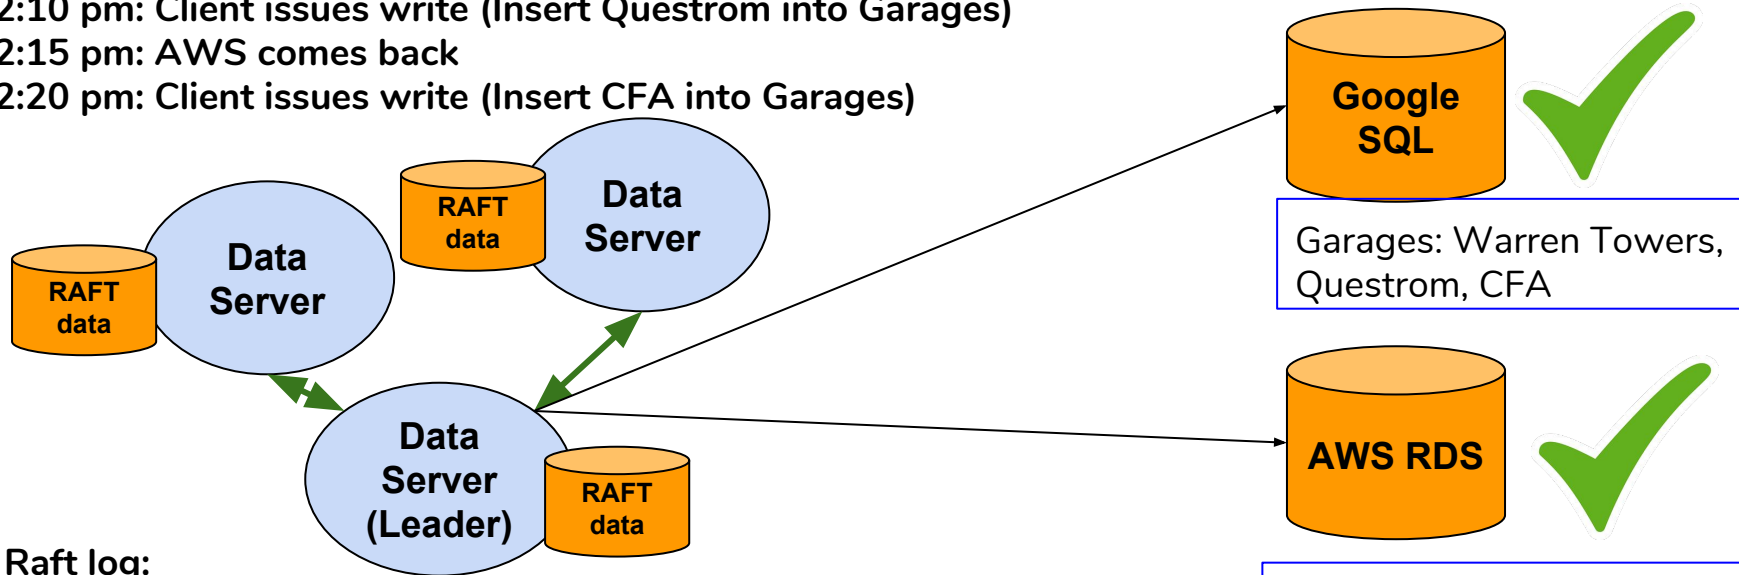
2:00 pm: Client issues write (Insert Warren Towers into Garages)

2:05 pm: AWS RDS goes down

2:10 pm: Client issues write (Insert Questrom into Garages)

2:15 pm: AWS comes back

2:20 pm: Client issues write (Insert CFA into Garages)



## Raft log:

{GCP\_DOWN, 0}

{AWS\_DOWN, **2:10pm**}

{Insert Warren Towers into Garages, timestamp: 2:00pm}

{Insert Questrom into Garages, **timestamp: 2:10pm**}

{Insert CFA into Garages, **timestamp: 2:20**}

# Database Recovery Example


2:00 pm: Client issues write (Insert Warren Towers into Garages)

2:05 pm: AWS RDS goes down

2:10 pm: Client issues write (Insert Questrom into Garages)

2:15 pm: AWS comes back

2:20 pm: Client issues write (Insert CFA into Garages)



## Raft log:

{GCP\_DOWN, 0}

{AWS\_DOWN, 0}

{Insert Warren Towers into Garages, timestamp: 2:00pm}

{Insert Questrom into Garages, timestamp: 2:10pm}


{Insert CFA into Garages, timestamp: 2:20}

# DNS A-NAME rotation

- What is DNS A-NAME rotation?
- DNS maps a domain name to an IP address
- We implement ANAME rotation by mapping single DNS name to multiple IP addresses
- The DNS service heartbeats with the hosts to ensure they are up before choosing the IP, so it knows if any of the machines is down
- Whenever `api.cloud-hydra.com` is requested, it uses round robin to either send it to AWS or GCP load balancers

# DEMO TIME!

# Sprint4 Burndown



# Next Steps

- Create a test suite to test Hydra
  - Unit tests
  - System and E2E tests
  - Metrics and timing characteristics
  - Validate consistency claims via test results
- DNS failover for application servers
- Identify behaviors if Hydra fails
- Enhance UI/UX of the frontend

**Thank you!**  
**Questions?**