

Cloud-Hydra: A Cloud Native Multi-Cloud Defensive Load Balancing Framework

Josh Stern, Rachid Tak Tak, Julian Trinh, Filip Vukelic

SPRINT 5

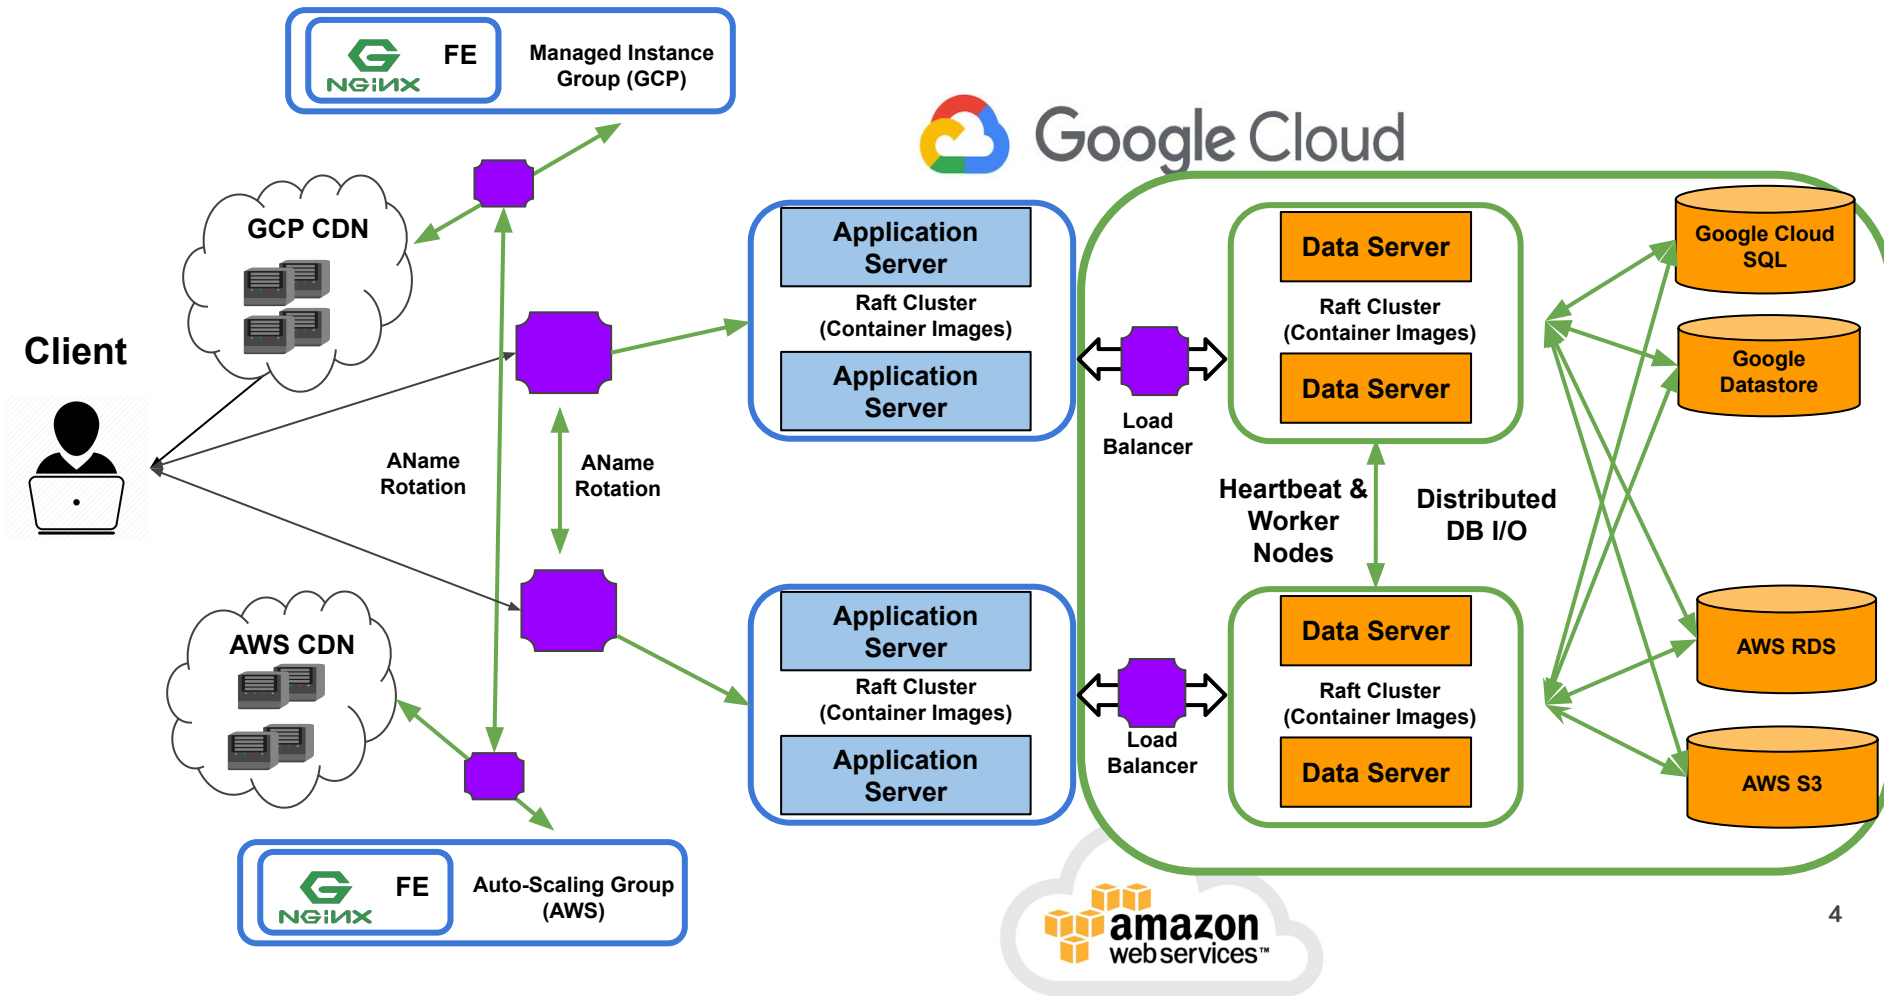


Sprint History

- **Sprint 1- Built fullstack web app in GCP and AWS**
 - Setup GCP and AWS infrastructure
- **Sprint 2 - Unified Data Layer**
 - Raft consensus
 - Cross-cloud consensus, single cloud load balancing
- **Sprint 3 - Cross cloud forwarding and recovery**
 - Leader Forwarding
 - Data layer Load Balancing
 - Database Recovery (single cloud)
- **Sprint 4 - Stitching together both clouds**
 - Synchronize cross cloud infra
 - Multi-Cloud DB Recovery
 - Stateless application layer

Goals of Sprint 5

- DNS load balancing
 - Front end
 - Application layer
- End-to-end testing of the entire system via UI
- Wallet Crushing: benchmarking throughput and latency



End-to-end Interconnection

- Remaking our infrastructure to new GCP and AWS emails and accounts
- Front end DNS ANAME rotation
- Since application layer is stateless, also applied DNS ANAME rotation
- Fully interconnecting the entire infrastructure

Cloud Hardware Specs



- GCP load balancer
 - 2 vCPUs (2GHz)
 - 7.5 GB RAM
- GCP Raft Nodes
 - 1 vCPU (2GHz)
 - 3 GB RAM
- GCP SQL
 - 1 vCPU
 - 3.75 GB



- AWS load balancer
 - 1 vCPU
 - 1 GB RAM
- AWS Raft Nodes
 - 1 vCPU
 - 1 GB RAM
- AWS RDS
 - 1 vCPU
 - 1 GB

Test Suite

- Script sends 200 requests in a loop
 - send next request only after previous one receives response
- Read payload 65 Bytes
- Write payload 36 Bytes

Latency - Hitting Multi-Cloud DBs vs Single

Latency(ms)	AWS	GCP	Both
Reads	47.03	39.29	44.41
Writes	73.65	68.19	76.92

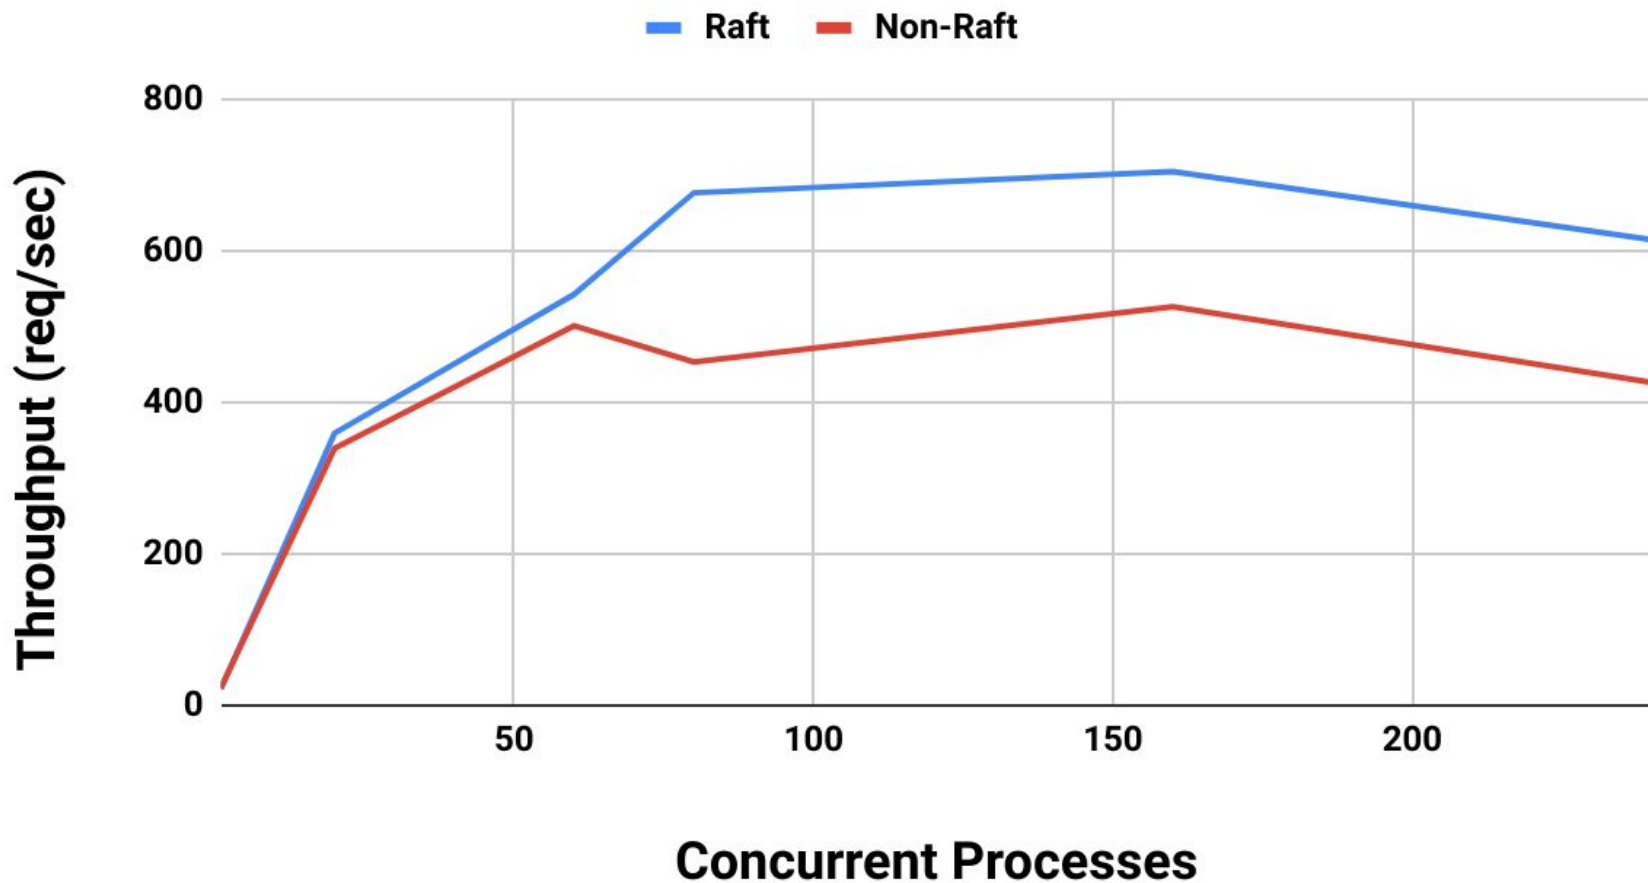
Read Throughput vs Number of Processes

Number of Concurrent Processes	Throughput (Requests / sec)	Average Latency (ms)
1	21.63	46.21
20	359	51.56
60	542	121.065
80	676	190.36
160	704	224.56
240	614	406.46

Read Throughput vs Number of Processes (no Raft Cluster)

Number of Concurrent Processes	Throughput (Requests / sec)	Average Latency (ms)
1	22.86	45.29
20	338.8	56.05
60	500.52	118.32
80	452.8	176.63
160	526	303.36
240	426	503.98

Reads: Raft and Non-Raft



Read Throughput vs Number of Mixed Processes

Number of Concurrent Processes	Throughput (Requests / sec)	Average Latency (ms)
30 Reads / 30 Writes	697.5	43.02
40 Reads / 40 Writes	714.28	56.05
80 Reads / 80 Writes	526.9	241.63
120 Reads / 120 Writes	276.38	457.07

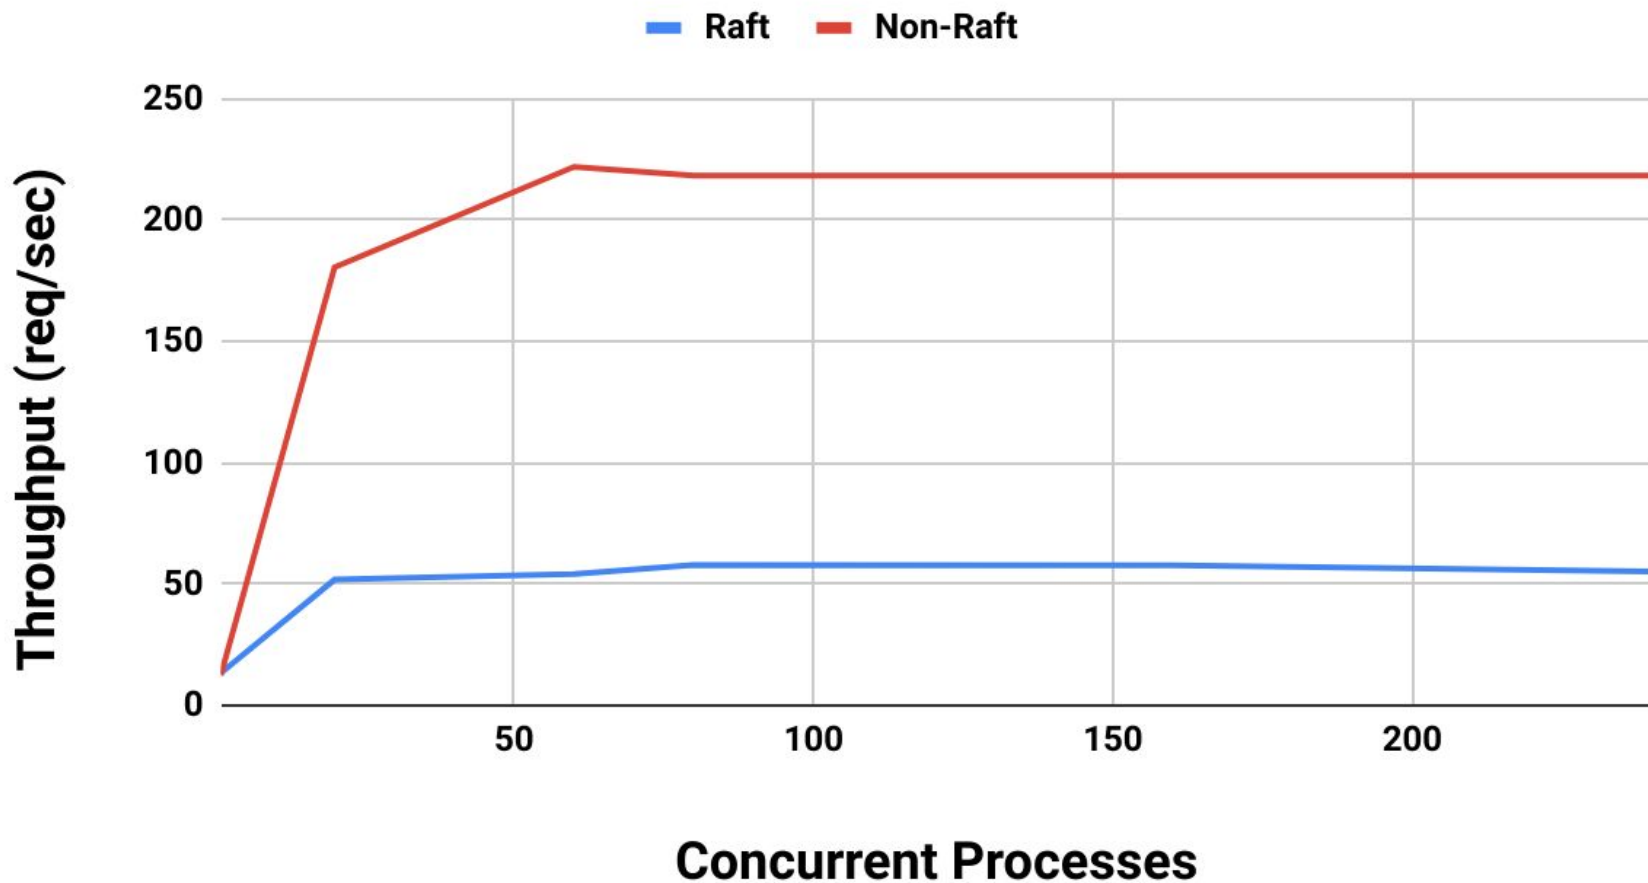
Write Throughput vs Number of Processes

Number of Concurrent Processes	Throughput (Requests / sec)	Average Latency (ms)
1	13.3	75.19
20	51.83	389
60	55.2	1087.36
80	57.92	1382.42
160	57.78	2768.75
240	55.2	4231

Write Throughput vs Number of Processes (no Raft Cluster)

Number of Concurrent Processes	Throughput (Requests / sec)	Average Latency (ms)
1	12.37	80.83
20	180.62	110.72
60	222	269.69
80	218.4	365.74
160	56% dropped requests	N/A
240	46% dropped requests	N/A

Writes: Raft and Non-Raft



Write Throughput vs Number of Mixed Processes


Number of Concurrent Processes	Throughput (Requests / sec)	Average Latency (ms)
30 Reads / 30 Writes	57.6	537.28
40 Reads / 40 Writes	54	739.86
80 Reads / 80 Writes	59.2	1397.14
120 Reads / 120 Writes	56.29	2318.2

Analysis

- Reads benefit from being in a load-balanced cluster
 - Requests can be shared across multiple-machines
 - Latency of the cluster was worse until ~160 concurrent connections
- Writes within a public cloud domain (Google -> Google) are faster than cross cloud due to locality
- Consensus and cluster-level locking reduce peak write throughput by 3.5x
 - Raft log is unbounded
 - The actual writes to the DB are small
 - Sensitive to network

DEMO TIME!

Sprint 5 Burndown



Challenges

- Needed to re-create entire infrastructure
- Throughput tests are highly network dependent
- Cached DNS-IP mapping

Future Optimizations

- Writes are slow because the leader has to lock the raft log for performing consensus and then db writes
 - Help?
- Bound the Raft log
 - Do we need to retain a log if both DBs are up?
 - Delete Raft entries after a certain size
- Use a more performant consensus library or algorithm
 - Ectd Raft (used by Kubernetes)
 - Leaderless Paxos

Thank you!

Questions?

